

# Final Report

## Chess State Prediction from a Video Stream

Calum Wallbridge

Submitted in accordance with the requirements for the degree of  
Computer Science

2021/22

<Module code and name>

The candidate confirms that the following have been submitted.

Items	Format	Recipient(s) and Date
Final Report	PDF file	Uploaded to Minerva (DD/MM/YY)
<Example> Scanned participant consent forms	PDF file / file archive	Uploaded to Minerva (DD/MM/YY)
<Example> Link to online code repository	URL	Sent to supervisor and assessor (DD/MM/YY)
<Example> User manuals	PDF file	Sent to client and supervisor (DD/MM/YY)

The candidate confirms that the work submitted is their own and the appropriate credit has been given where reference has been made to the work of others.

I understand that failure to attribute material which is obtained from another source may be considered as plagiarism.

(Signature of Student) \_\_\_\_\_

## Summary

<Concise statement of the problem you intended to solve and main achievements (no more than one A4 page)>

Immediately explain what you've added compared to other implementations.

## **Acknowledgements**

<The page should contain any acknowledgements to those who have assisted with your work. Where you have worked as part of a team, you should, where appropriate, reference to any contribution made by other to the project.>

Note that it is not acceptable to solicit assistance on ‘proof reading’ which is defined as the “the systematic checking and identification of errors in spelling, punctuation, grammar and sentence construction, formatting and layout in the test”; see

[https://www.leeds.ac.uk/secretariat/documents/proof\\_reading\\_policy.pdf](https://www.leeds.ac.uk/secretariat/documents/proof_reading_policy.pdf)

# Contents

<b>1</b>	<b>Introduction and Background Research</b>	<b>2</b>
1.1	Introduction . . . . .	2
1.2	Literature review . . . . .	2
1.2.1	A Short History of Computer Vision . . . . .	2
1.2.2	Tools . . . . .	3
1.2.3	Computer Vision for Chess . . . . .	3
1.2.4	Prior Work From the Author . . . . .	3
<b>2</b>	<b>Methods</b>	<b>4</b>
2.1	Data Collection . . . . .	4
2.1.1	Sensors . . . . .	4
2.1.2	Auto-Labeling . . . . .	4
2.1.3	Dataset Versioning . . . . .	4
2.2	Model Training . . . . .	5
2.2.1	Architectures . . . . .	5
2.2.2	Transfer Learning . . . . .	6
2.2.3	Augmentation . . . . .	6
2.2.4	Experiment Tracking . . . . .	6
2.3	Inference . . . . .	6
2.3.1	The Reliable Approach . . . . .	6
2.3.2	The Not-So-Reliable Approach . . . . .	6
<b>3</b>	<b>Results</b>	<b>7</b>
3.1	Model Evaluation . . . . .	7
3.1.1	Board Identification . . . . .	7
3.1.2	Multitask Learning . . . . .	7
3.1.3	Piece Recognition . . . . .	7
3.1.4	Piece Detection . . . . .	7
3.1.5	Deep Dive into CNNs . . . . .	7
3.2	Realtime Analysis . . . . .	7
3.2.1	Trails . . . . .	8
3.3	Comparison to Existing Solutions . . . . .	8
<b>4</b>	<b>Discussion</b>	<b>9</b>
4.1	Conclusions . . . . .	9
4.2	Ideas for future work . . . . .	9
	<b>References</b>	<b>10</b>
	<b>Appendices</b>	<b>11</b>

<i>CONTENTS</i>	1
<b>A Self-appraisal</b>	<b>11</b>
A.1 Critical self-evaluation . . . . .	11
A.2 Personal reflection and lessons learned . . . . .	11
A.3 Legal, social, ethical and professional issues . . . . .	11
A.3.1 Legal issues . . . . .	11
A.3.2 Social issues . . . . .	11
A.3.3 Ethical issues . . . . .	11
A.3.4 Professional issues . . . . .	11
<b>B External Material</b>	<b>12</b>

# Chapter 1

## Introduction and Background Research

### 1.1 Introduction

Algorithms such as Deep Blue [1], AlphaZero [2] and more recently Player of Games [3] have enabled computers to out smart the smartest humans at the game of Chess. But all these algorithms are bound to the digital world, rendered useless when competing against humans on a real board. This project aims to explore a major component of this: vision. Consider the vision problem for chess to be two-fold: what is the current board state and where are all of the pieces? With this information, in combination with the previous algorithms and a robot arm, the computer is no longer bound to the digital world. In particular this project's will focus on the former, that is, to produce and present a solution for determining the state of a chess board from a video stream. A solution reliable enough to live up to the likes of AlphaZero in a robotic solution. There will be a focus on deep learning techniques, with consideration for best practice in machine learning operations. There is also the aim to share the tools to more easily manage and create new datasets in this area of chessboard state recognition. Something called for by [4] as a serious challenge and priority for future research.

### 1.2 Literature review

For humans the hard part of chess is planning, this is not the case for computers, instead recognition and localisation of objects in 3D space, along with manipulation, present much greater challenges. [5]

Make reference to humans huge allocation of resources to vision. [6] Why is it so hard for computers then? It's an inverse problem. Compare to solving the decision problem (minimax). The statistical calculation of whether to trade Queen's or block with a pawn has now become trivial.

#### 1.2.1 A Short History of Computer Vision

##### Classical Techniques

How we can use classical techniques to understand images. Neural networks have taken over almost all of the heavy lifting for high level inference. Give an indication of time scale here.

##### Image Recognition

the way back in 1980 [7] convolutions showed promise in simple computer vision tasks, convolutions since have showed extreme promise [8] Le Ye Cunn's work with convolutions [9] and MNIST [10] and more recently ImageNet [11] having become incredibly well known. New methods such as Transformers from the world have NLP have generalised the convolution operation have

proved very successful and lots of work here has been applied to vision. [1] [2] (Attention is all you need, ViT, generic model from deepmind)

### Object Detection

But in most applications there will be many things we want to recognise in an image. The RCNN [3] enters. How Faster-RCNN improves on this [4]. Why YOLO [5] has been so successful (realtime). Transformers are applicable here too.

### Instance Segmentation

Why bounding boxes are not enough. What is segmentation? Why instance segmentation is what we actually want. [6]

### Adding More Dimensions

The real world is not perceived in static 2d images. How do we add an understanding of 3 dimensions and time in our computer vision models? [7] Important for localisation in the real world. Important for understanding things like object permanence.

## 1.2.2 Tools

Dataset Collection??

Chess Engine??

Deep Learning Library??

## 1.2.3 Computer Vision for Chess

Despite chess being a very narrow application of computer vision, the amount of research effort gone into the problem is not insignificant, in contrary to some papers claims [8]. A variety of approaches have been tested and explored for which the following section will attempt to fairly summarise. Finding the chess board is the first challenge, the most common solution as in [9] is to use the ‘findChessboard’ function in openCV for camera calibration. This \*\*how does this work\*\*.

## 1.2.4 Prior Work From the Author

Mention robotic arm for two counter board games and automatic differentiation library.



# Chapter 2

## Methods

<Everything that comes under the ‘Methods’ criterion in the mark scheme should be described in one, or possibly more than one, chapter(s).>

### 2.1 Data Collection

At the heart of any machine learning project is the data. It is as important, often more important, than the code and presents many interesting challenges. **\*\*Why is this?\*\*** Discussed in the following sections are some of the challenges and decisions that were considered.

#### 2.1.1 Sensors

The eminent challenge is acquiring data in the first place. This is highly context dependant, but as vision is primarily focused on spatial awareness the discussion will be limited to the sensors that can measure it.

Sensor choice is an important choice for any robotics application as there are important tradeoffs, as with any engineering challenge, which must be considered.

**\*\*Outline some of the tradeoffs between spatial sensors\*\***

One important distinction to make is the difference between training and inference.

Requirements at the time of training may differ significantly to the requirements at inference.

Processing power, energy supply and realtime operation are some of the constraints that will have to be met when considering different sensors.

The sensor used throughout this project is the RealSense SR305 which is a RGB-D camera using structured and coded light to determine depth, it functions best indoors or in a controlled lighting situation. For the reasons outlined above the RGB camera stream is mainly relied upon but there will be some discussion and comparison of piece detection with the depth sensor.

#### 2.1.2 Auto-Labelling

A closely related challenge of acquiring the data is that of labelling it too. It is widely known that neural networks scale with the number of examples []. This will be explicitly explored for Chess Vision in section 3.1. This however poses the question about how do we get access to a lot of labelled data for chess. **\*\*Some examples of other auto labelling techniques\*\*** The overall approach I took...

#### 2.1.3 Dataset Versioning

With all this data the next challenge becomes self evident. It is concerned with the question: How do we manage all of this? Some of the problems... and why you need versioning...

Transitioning from git lfs to aws s3. Perhaps a quick mention of other solutions. How the Game class solves some of these challenges for us.

## 2.2 Model Training

Classification is all that's needed.

### 2.2.1 Architectures

To start, a good baseline is found. A good baseline is a simple model to understand and easy to get decent results with. For classification, the pathological baseline could be a random model, which could easily be extended to be weighted by count of examples for each class in the training set. [Add mathematical formulation for demonstration (multinomial)]

This is common practice in exploratory machine learning as it keeps the research from getting lost, so that the transitions made are always from a known and working state. It becomes very easy to see if an experiment is not working and easy to go back.

Keeping to this strategy, the multilayer perceptron (MLP) or fully connected network was the first neural network to be explored. By starting with the MLP, all complexity from the network is stripped away so the more extraneous elements such as the training loop and evaluation metrics can be built and tested. [perhaps mention purposefully overfitting]

Once the full training and evaluation structure is functional, new features can be incrementally added and architectures explored.

As mentioned in Literature review convolution operations, and in particular differentiable convolutional operations have had monumental impact on the field of computer vision and so this was the first experiment.

Quite quickly, especially with a limited dataset, overfitting becomes a major problem to overcome and so many experiments were positioned To solve this problem. Pooling, dropout layers and skip connections are among some of these.

ConvNext <https://arxiv.org/pdf/2201.03545.pdf>

### Optimizer

An important part of any neural network training loop is the optimiser and so we will explore the effect of swapping these out.

[put into results] AdamW was found produce more attractive results from comparing SGD with/without momentum and Adam and AdamW Learning Rate, Weight Regularization.

### Batch Normalization

<https://arxiv.org/abs/1502.03167> and renorm for small batch sizes

<https://arxiv.org/abs/1702.03275>

### Dropouts

<https://arxiv.org/abs/1207.0580> shouldn't be used after convolution layers there has been some work <https://arxiv.org/abs/1904.03392>

Pooling and DropPaths

### 2.2.2 Transfer Learning

There appears to be a trend occurring in the deep learning space. Some organisation spends millions training an impossibly large neural network and others more and more are using these models, often fine-tuning for their own use cases. [1] uses these large models as fixed feature extractors.

This approach makes sense as it is impractical to retrain huge neural networks that take weeks, millions of dollars and wasteful amounts of energy to train [1].

In the case of CNNs we can see the features that kernels in the early layers learn [1] are often very simple shapes and will be common for all computer vision tasks. This will be explored further in the results section as we visualise kernels from both random initialised models and pretrained models.

### 2.2.3 Augmentation

Data augmentation is a strategy every machine learning practitioner wants to have in their arsenal. It addresses the data problem, allows us to truly leverage the data we have and generalise our models further. The MNIST [1] dataset itself was created using data augmentation. In the case of chess piece classification, the correct augmentations or transformations should be chosen. Go on to list the transformations used and why.

### 2.2.4 Experiment Tracking

Express importance of experiment tracking. Some of the solutions [1] and their tradeoffs. Why a guild was chosen and how it was used.

## 2.3 Inference

### 2.3.1 The Reliable Approach

As in [1] the most reliable way to implement inference during play is for a user to input when a move has been complete. This is not ideal because...

### 2.3.2 The Not-So-Reliable Approach

What if the system constantly updates with each new frame? Why this is what we want, but is not as important for finding board state. Some of the challenges involved and how we overcame them.

#### Detecting Motion

Difference between 2 frames. SIFT / SURF

#### Board State and Vision State

Leveraging the chess engine to give more reliable results for robot play or pgn recording.

# Chapter 3

## Results

<Results, evaluation (including user evaluation) *etc.* should be described in one or more chapters. See the ‘Results and Discussion’ criterion in the mark scheme for the sorts of material that may be included here.>

### 3.1 Model Evaluation

#### 3.1.1 Board Identification

Aruco vs Chessboard vs heatmap.

Compare results of 5 examples. Why chessboard [] will struggle with peices on the board. Breif intro to Aruco. For capturing data Aruco is very reliable, but won’t always be possible to add onto the board. Why heatmap solution looks like the best appraoch in the realword with unknown environments. []

Time comparision for the 3. And compare accuracy of heatmap against aruco.

#### 3.1.2 Multitask Learning

How does it compare?

#### 3.1.3 Piece Recognition

Use top-1 and top-5 to measure performance of different models. Will include basic evaluation. What happens what you increase layers, use more data, data augmentation. Include Recall / Specificity / Sensitivity

#### 3.1.4 Piece Detection

A benefit of having the depth sensor is an easier way to detech piece presence. Fixed threshold vs clustering. Adding a margin. How we actually did it. Using paried T-test to evaulate. Using a neural network instead as an additional class with our piece recognition network. Compare that two having a two stage network, the first for piece detection and the second for recognition.

#### 3.1.5 Deep Dive into CNNs

Visualising convolutional layers to analyse effectiveness.

### 3.2 Realtime Analysis

Frames per second. Problems.

**3.2.1 Trails**

table of end-to-end experiments

**3.3 Comparison to Existing Solutions**

Summaries how my solution compares to others. Need a table?

# Chapter 4

## Discussion

<Everything that comes under the ‘Results and Discussion’ criterion in the mark scheme that has not been addressed in an earlier chapter should be included in this final chapter. The following section headings are suggestions only.>

### 4.1 Conclusions

### 4.2 Ideas for future work

Firstly it would be nice for more extreme camera angles to work just as reliably. If assumptions about the environment are to be kept minimal then localising pieces in 3 dimensional space should be a requirement for robotic manipulation to be possible.

It is the author’s belief that "3d reconstruction" would kill two birds with one stone and the direction that future research should focus.

Explain. Similar work in other areas.

# References

- [1] D. Parikh, N. Ahmed, and S. Stearns. An adaptive lattice algorithm for recursive filters. *Acoustics, Speech and Signal Processing, IEEE Transactions on*, 28(1):110–111, 1980.

# Appendix A

## Self-appraisal

<This appendix should contain everything covered by the 'self-appraisal' criterion in the mark scheme. Although there is no length limit for this section, 2—4 pages will normally be sufficient. The format of this section is not prescribed, but you may like to organise your discussion into the following sections and subsections.>

### A.1 Critical self-evaluation

### A.2 Personal reflection and lessons learned

Surprised at how effective transfer learning is and the significance of it's place in the future. The importance of experiment tracking for research.

### A.3 Legal, social, ethical and professional issues

<Refer to each of these issues in turn. If one or more is not relevant to your project, you should still explain *why* you think it was not relevant.>

#### A.3.1 Legal issues

#### A.3.2 Social issues

#### A.3.3 Ethical issues

#### A.3.4 Professional issues



# Appendix B

## External Material

<This appendix should provide a brief record of materials used in the solution that are not the student's own work. Such materials might be pieces of codes made available from a research group/company or from the internet, datasets prepared by external users or any preliminary materials/drafts/notes provided by a supervisor. It should be clear what was used as ready-made components and what was developed as part of the project. This appendix should be included even if no external materials were used, in which case a statement to that effect is all that is required.>